

Tabular-only vs. tabular+BERT accuracy across the 2×2 study grid
(dashed line = chance level; significance = paired bootstrap test, $\alpha = 0.05$, $n = 2000$ resamples)

